

# **Chapter 9**

## **Nonblocking I/O**



### **Content:**

1. An Example Client and Server
2. Buffers
  - 2.1 Creating Buffers
  - 2.2 Filling and Draining
  - 2.3 Bulk Methods
  - 2.4 Data Conversion
  - 2.5 View Buffers
  - 2.6 Compacting Buffer
  - 2.7 Duplicating Buffers
  - 2.8 Slicing Buffers
  - 2.9 Marking and Resetting
  - 2.10 Object Methods
3. Channels
  - 3.1 SocketChannel
  - 3.2 ServerSocketChannel
  - 3.3 The Channels Class
  - 3.4 Asynchronous Channels
  - 3.5 Socket Options
4. Readiness Selection
  - 4.1 The Selector Class
  - 4.2 The SelectionKey Class

# Java I/O:

## Two typical I/O models

- Stream-oriented I/O
  - Movement of single bytes, one at a time
  - Byte streams and character streams
  - Simple
- Block-oriented I/O
  - Dealing with data in blocks, especially for bulk data transfers
    - A low-level data transfer mechanism
  - Channels and buffers
  - Faster



## Java I/O

- Original I/O package: java.io.\*
- New I/O package: NIO (JDK 1.4+)

- **Channels and Buffers**


- Data is always read from a channel into a buffer, or written from a buffer to a channel

- **Non-blocking I/O**

- After asking a channel to read data into a buffer, a thread can do something else while the channel reads data into the buffer

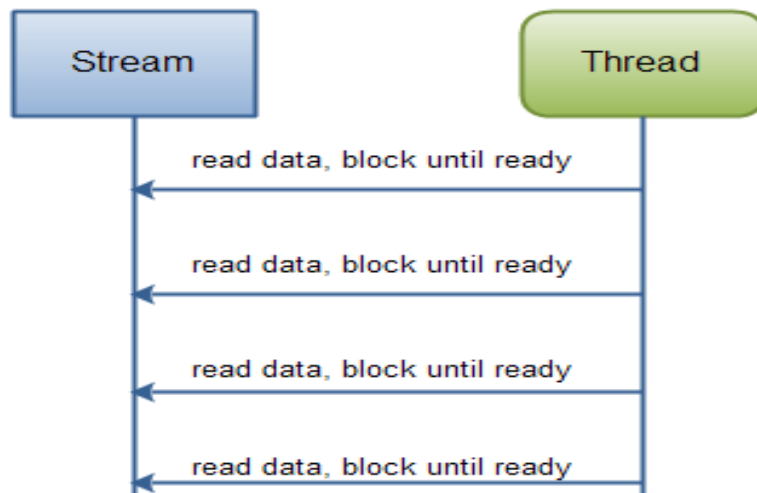
- **Selector:** an object that can monitor multiple channels for events

- A thread can monitor multiple channels for data

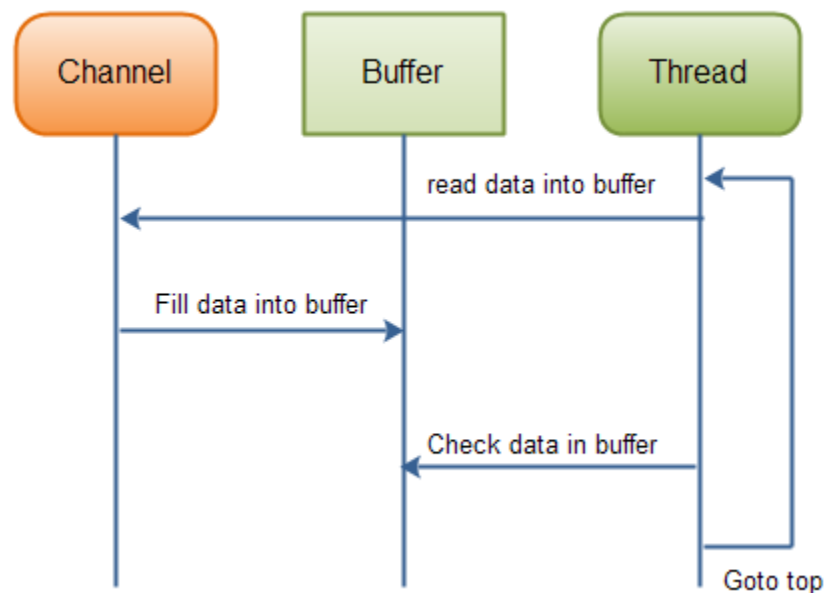


# Java IO vs NIO:


## Java IO: Reading data form a blocking stream




## Java NIO: Reading data from a channel until all needed data is in buffer



Java IO: A classic IO server design one connection handled by one thread



Java NIO: A single thread managing multiple connections



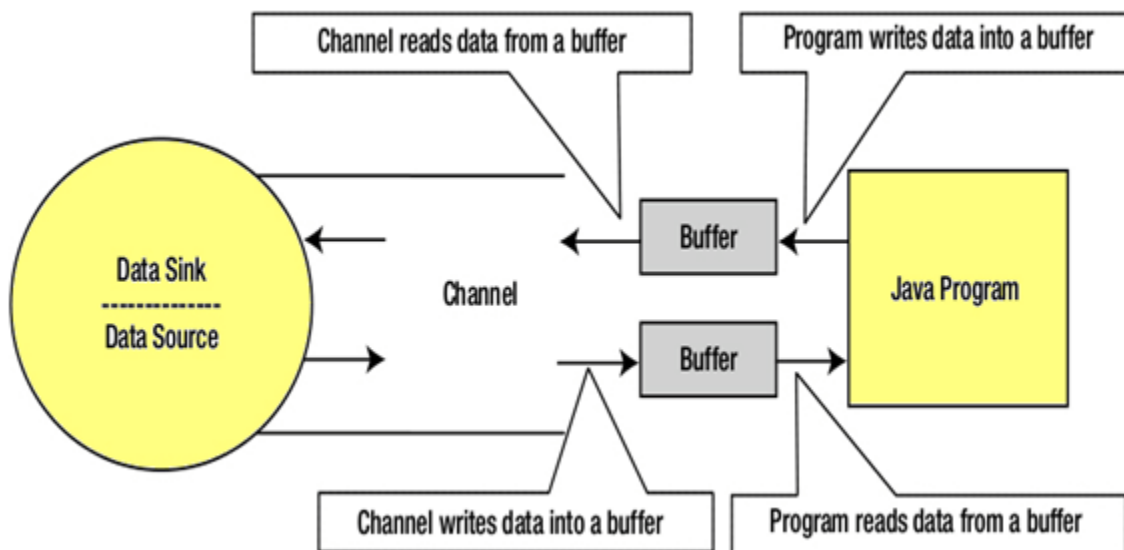
Java NIO's **selectors** allow a single thread to monitor multiple channels of input

IO	NIO
It is based on the Blocking I/O operation	It is based on the Non-blocking I/O operation
It is Stream-oriented	It is Buffer-oriented
Channels are not available	Channels are available for Non-blocking I/O operation
Selectors are not available	Selectors are available for Non-blocking I/O operation

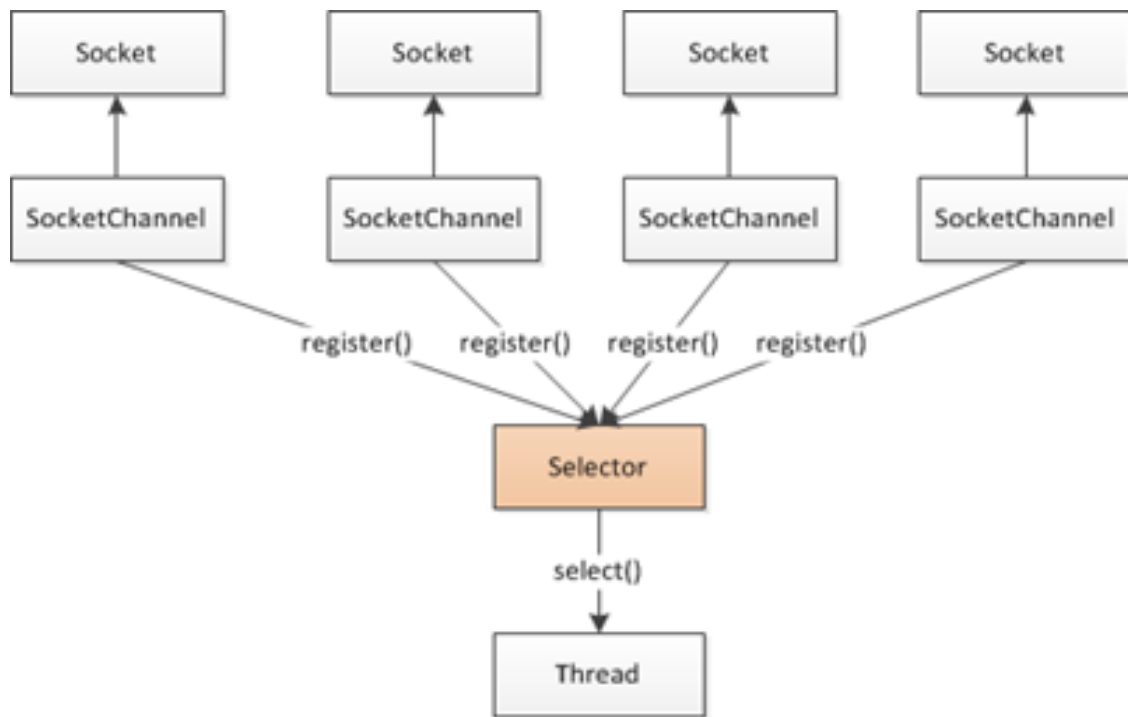
## java.nio Package:

- Central abstractions of the NIO APIs
  - **Buffers:** containers for for a fixed amount of data of a specific primitive type
    - ByteBuffer (MappedByteBuffer), CharBuffer
    - ShortBuffer, IntBuffer, LongBuffer
    - FloatBuffer, DoubleBuffer
  - **Channels:** represent connections to entities capable of performing I/O operations

- FileChannel, DatagramChannel, SocketChannel, ServerSocketChannel
- **Selectors** and selection keys, which together with selectable channels: define a multiplexed, non-blocking I/O facility
- **Charsets** and their associated decoders and encoders: translate between bytes and Unicode characters



*Interaction between channel, buffers, java program, data source and data sink*



## Buffers:

- Essential properties of a buffer
  - **Capacity**: the number of elements it contains
    - Specified when the Buffer is constructed and cannot be changed (similar to an array)
    - Never changes
  - **Limit**: specifies the current occupancy (valid data in the range of 0 to limit-1)
    - Never greater than its capacity
  - **Position**: the index of the next element to be read or written



-Never greater than its limit

## Mark:


A client-specified index in the buffer. It is set at the current position by invoking the `mark()` method. The current position is set to the marked position by invoking `reset()`:

*public final Buffer mark()*

*public final Buffer reset()*

If the position is set below an existing mark, the mark is discarded

Example:



**flip()** : set Position to 0 and Limit to last Position

**clear()**: set Position to 0 and Limit to Capacity

**rewind()**:set Position to 0 only

There are two methods that return information about the buffer but don't change it:

**The remaining()** method returns the number of elements in the buffer between the current position and the limit.

**The hasRemaining()** method returns true if the number of remaining elements is greater than zero:

```
public final int remaining()
```

```
public final boolean hasRemaining()
```

## Creating Buffers:

**Empty buffers** are normally created by **allocate methods**.

Buffers that are prefilled with data are created by wrap methods.

The **allocate methods** are often useful for **input**, and the **wrap methods** are normally used for **output**

**Allocate methods are two type:**

### 1.1 Allocation:

The basic `allocate()` method simply returns a new, empty buffer with a specified fixed capacity.

For example, these lines create byte and int buffers, each with a size of 100:

```
ByteBuffer buffer1 = ByteBuffer.allocate(100);
```

```
IntBuffer buffer2 = IntBuffer.allocate(100);
```

you could read a large chunk of data into a buffer using a channel and then retrieve the array from the buffer to pass to other methods:

```
byte[] data1 = buffer1.array();
```

```
int[] data2 = buffer2.array();
```

## **1.2 Direct Allocation:**

The ByteBuffer class (but not the other buffer classes) has an additional `allocateDirect()` method that may not create a backing array for the buffer.

The VM may implement a directly allocated ByteBuffer using direct memory access to the buffer on an Ethernet card, kernel memory, or something else

From an API perspective, `allocateDirect()` is used exactly like `allocate()`:

```
ByteBuffer buffer = ByteBuffer.allocateDirect(100);
```

## Wrapping:

If you already have an array of data that you want to output, you'll normally wrap a buffer around it, rather than allocating a new buffer and copying its components into the buffer one at a time. For example:

```
byte[] data = "Some data".getBytes("UTF-8");
```

```
ByteBuffer buffer1 = ByteBuffer.wrap(data);
```

```
char[] text = "Some text".toArray();
```

```
CharBuffer buffer2 = CharBuffer.wrap(text);
```

## Filling and Draining the buffer:

- Buffers are designed for sequential access.
- Recall that each buffer has a current position identified by the `position()` method that is somewhere between zero and the number of elements in the buffer, inclusive.
- The buffer's position is incremented by one when an element is read from or written to the buffer.

For example, suppose you allocate a `CharBuffer` with capacity 12, and fill it by putting five characters into it:

```
CharBuffer buffer = CharBuffer.allocate(12);  
buffer.put('H');  
buffer.put('e');  
buffer.put('l');  
buffer.put('l');  
buffer.put('o');
```

The position of the buffer is now 5. This is called filling the buffer

```
ByteBuffer buffer = ByteBuffer.allocate(512);  
String str = "Hello";  
byte[] data = str.getBytes();  
buffer.put(data);
```

- You can only fill the buffer up to its capacity.
- If you tried to fill it past its initially set capacity, the `put()` method would throw a `BufferOverflowException`.
- If you now tried to `get()` from the buffer, you'd get the null character (`\u0000`) that Java initializes char buffers with that's found at position 5. Before you can read the data you wrote in out again, you need to flip the buffer:

```
buffer.flip();  
int limit = buffer.limit();  
byte[] data = new byte[limit];  
buffer.get(data);  
System.out.println(new String(data));
```

## **Bulk Methods:**

- it's often faster to work with blocks of data rather than filling and draining one element at a time.
- The different buffer classes have bulk methods that fill and drain an array of their element type.
- For example, ByteBuffer has put() and get() methods that fill and drain a ByteBuffer from a preexisting byte array or subarray:

```
public ByteBuffer get(byte[] dst, int offset, int length)  
public ByteBuffer get(byte[] dst)
```

```
public ByteBuffer put(byte[] array, int offset, int length)  
public ByteBuffer put(byte[] array)
```

## View Buffers:

View buffers are useful when you want to access a portion of a buffer's data without modifying it, or when you want to share the same data between multiple buffers.

When you create a view buffer, it shares the same data as the original buffer, but has its own position, limit, and mark values.

View Buffers are created with one of these six methods in `ByteBuffer`:

***public abstract ShortBuffer asShortBuffer()***

***public abstract CharBuffer asCharBuffer()***

***public abstract IntBuffer asIntBuffer()***

***public abstract LongBuffer asLongBuffer()***

***public abstract FloatBuffer asFloatBuffer()***

***public abstract DoubleBuffer asDoubleBuffer()***

## Compacting Buffers:

In Java, **`ByteBuffer.compact()`** is a method used to compact the buffer, which means it moves any remaining data to the beginning of the buffer and makes room for new data to be written.

## Here's how it works:

If there is no remaining data in the buffer, then it simply **clears the buffer and resets its position and limit** to their initial values.

If there is **remaining data in the buffer**, it copies this data to the **beginning of the buffer, starting from position 0**.

It then sets the position to the number of bytes copied and the limit to the buffer's capacity.

The compact method is useful when we want to **reuse** a buffer without having to create a new one.

By compacting the buffer, we can ensure that there is enough room for new data to be written without overwriting the remaining data in the buffer.

Most writable buffers support a compact() method:

***public abstract ByteBuffer compact()***

***public abstract IntBuffer compact()***

***public abstract ShortBuffer compact()***

***public abstract FloatBuffer compact()***

***public abstract CharBuffer compact()***

***public abstract DoubleBuffer compact()***



## Duplicating Buffers:

It's often desirable to make a copy of a buffer to deliver the same information to two or more channels.

The `duplicate()` methods in each of the six typed buffer classes do this:

***public abstract ByteBuffer duplicate()***

***public abstract IntBuffer duplicate()***

***public abstract ShortBuffer duplicate()***

***public abstract FloatBuffer duplicate()***

***public abstract CharBuffer duplicate()***

***public abstract DoubleBuffer duplicate()***

The return values are not clones. The duplicated buffers share the same data, including the same backing array if the buffer is indirect. **Changes to the data in one buffer are reflected in the other buffer.**

## Slicing Buffers:

- Slicing a buffer is a slight variant of duplicating. Slicing also creates a new buffer that shares data with the old buffer

- This new buffer starts from the position of the original buffer and extends to its limit.
- The capacity of the new buffer is equal to the remaining capacity of the original buffer (limit - position).
- Changes made to the data of the new buffer are reflected in the original buffer and vice versa.

There are separate `slice()` methods in each of the six typed buffer classes:

*`public abstract ByteBuffer slice()`*

*`public abstract IntBuffer slice()`*

*`public abstract ShortBuffer slice()`*

*`public abstract FloatBuffer slice()`*

*`public abstract CharBuffer slice()`*

*`public abstract DoubleBuffer slice()`*

## Marking and Resetting:

The `mark()` method sets the **current position in the buffer as the "mark"**. You can call this method at any point in time.

The `reset()` method sets the **current position in the buffer to the previously set mark**, or to **zero** if a mark was never set.

*public final Buffer mark()*  
*public final Buffer reset()*

## Object Method:

Every class extends the Object class, which means that each class inherits certain methods from the Object class. These methods can be used to interact with objects of any class. The buffer classes all provide the usual equals(), hashCode(), and toString() methods.

- **equals()**: This method **compares two objects** for **equality** and returns true if the objects are equal, and false otherwise.
- **hashCode()**: This method returns a **hash code value** for the object. It is typically used in hashing algorithms and data structures.
- **toString()**: This method returns a **string representation** of the object. By default, it returns the class name and a hexadecimal representation of the object's hash code.

## Data Conversion:

- All data in Java ultimately resolves to bytes.
- Any primitive data type int, double, float, etc. can be written as bytes.
- For example, any sequence of four bytes corresponds to an int or a float (actually both, depending on how you want to read it).
- A sequence of eight bytes corresponds to a long or a double
- The ByteBuffer class (and only the ByteBuffer class) provides relative and absolute put methods that fill a buffer with the bytes corresponding to an argument of primitive type (except boolean) and relative and absolute get methods that read the appropriate number of bytes to form a new primitive datum.

***public abstract char getChar()***

***public abstract ByteBuffer putChar(char value)***

***public abstract char getChar(int index)***

***public abstract ByteBuffer putChar(int index, char value)***

***public abstract short getShort()***

***public abstract ByteBuffer putShort(short value)***

***public abstract short getShort(int index)***

***public abstract ByteBuffer putShort(int index, short value)***

***public abstract int getInt()***

***public abstract ByteBuffer putInt(int value)***

***public abstract int getInt(int index)***

***public abstract ByteBuffer putInt(int index, int value)***

***public abstract long getLong()***

***public abstract ByteBuffer putLong(long value)***

***public abstract long getLong(int index)***

***public abstract ByteBuffer putLong(int index, long value)***

***public abstract float getFloat()***

***public abstract ByteBuffer putFloat(float value)***

***public abstract float getFloat(int index)***

***public abstract ByteBuffer putFloat(int index, float value)***

***public abstract double getDouble()***

***public abstract ByteBuffer putDouble(double value)***

***public abstract double getDouble(int index)***

***public abstract ByteBuffer putDouble(int index, double value)***

## Byte buffer:

A byte buffer is a **data structure that holds a sequence of bytes**.

It provides a way to **manipulate and transfer** blocks of **binary data** efficiently.

Byte buffers are commonly used in network communication, file I/O, and other data-intensive applications where data needs to be transferred between different systems or processes.

They are also used in programming languages such as Java, C++, and Python to perform low-level operations on binary data.

Byte buffers are typically implemented as arrays of bytes with methods for reading and writing data.

They can be used for tasks such as converting data between different encodings, performing cryptographic operations, and manipulating binary data structures.

## An Example Client and Server:

A **Chargen server** is a software application or service that implements the Character Generator Protocol (Chargen), which is a simple protocol that generates a stream of characters and sends them to a remote host.

The Chargen server generates a continuous stream of characters, typically starting with **ASCII character 33 ("!")**, and continues until it reaches **ASCII character 126 ("~")**, at which point it starts over again.

## Example Chargen Client:

```
import java.io.IOException;  
import java.net.InetSocketAddress;  
import java.nio.ByteBuffer;  
import java.nio.channels.SocketChannel;  
import java.nio.charset.StandardCharsets;  
  
public class ChargenClient {  
    public static void main(String[] args) throws IOException {  
        // Connect to the server  
        SocketChannel socket = SocketChannel.open();  
        socket.connect(new InetSocketAddress("localhost", 8888));
```

```

// Read data from the server
ByteBuffer buffer = ByteBuffer.allocate(74);
while (socket.read(buffer) != -1) {
    buffer.flip();
    String data =
StandardCharsets.US_ASCII.decode(buffer).toString();
    System.out.print(data);
    buffer.clear();
}

// Close the connection
// socket.close();
}
}

```

## Example Chargen Server:

```

import java.io.IOException;
import java.net.*;
import java.nio.ByteBuffer;

```



```
import java.nio.channels.ServerSocketChannel;  
import java.nio.channels.Selector;  
import java.nio.channels.SelectionKey;  
import java.nio.channels.SocketChannel;  
import java.util.*;
```

```
public class Server2 {
```

```
    public static void main(String[] args) {
```

```
        byte[] rotation = new byte[95 * 2];
```

```
        for (byte i = ' '; i <= '~'; i++) {
```

```
            rotation[i - ' '] = i;
```

```
            rotation[i + 95 - ' '] = i;
```

```
        }
```

```
        ServerSocketChannel serverChannel;
```

```
        Selector selector;
```

```
        try {
```

```
            serverChannel = ServerSocketChannel.open();
```

```
            ServerSocket ss = serverChannel.socket();
```

```
            InetSocketAddress address = new  
InetSocketAddress(8888);
```

```

    ss.bind(address);

    serverChannel.configureBlocking(false);
    selector = Selector.open();
    serverChannel.register(selector,
SelectionKey.OP_ACCEPT);
    } catch (IOException ex) {
        ex.printStackTrace();
        return;
    }
    while (true) {
        try {
            selector.select();
        } catch (IOException ex) {
            ex.printStackTrace();
            break;
        }

        Iterator<SelectionKey>                iterator                =
selector.selectedKeys().iterator();

        while (iterator.hasNext()) {
            SelectionKey key = iterator.next();
            iterator.remove();

```

```

    try {
        if (key.isAcceptable()) {
            ServerSocketChannel server =
(ServerSocketChannel) key.channel();
            SocketChannel client = server.accept();
            System.out.println("Accepted connection from " +
client);
            client.configureBlocking(false);
            SelectionKey key2 = client.register(selector,
SelectionKey.OP_WRITE);
            ByteBuffer buffer = ByteBuffer.allocate(74);
            buffer.put(rotation, 0, 72);
            buffer.put((byte) '\r');
            buffer.put((byte) '\n');
            buffer.flip();
            key2.attach(buffer);
        } else if (key.isWritable()) {
            SocketChannel client = (SocketChannel)
key.channel();
            ByteBuffer buffer = (ByteBuffer) key.attachment();
            client.write(buffer);
            if (!buffer.hasRemaining()) {

```

```

        buffer.rewind();
        int first = buffer.get();
        buffer.rewind();
        int position = first - ' ' + 1;
        buffer.put(rotation, position, 72);
        buffer.put((byte) '\r');
        buffer.put((byte) '\n');
        buffer.flip();
    }

    client.register(selector,
SelectionKey.OP_WRITE, buffer);
}

} catch (IOException ex) {
    key.cancel();

    try {
        key.channel().close();
    } catch (IOException cex) {
    }

}

}

}

```

```
}  
  
}
```

## Channels:

Channels move blocks of data into and out of buffers to and from various I/O sources such as files, sockets, datagrams, and so forth.

### 1. SocketChannel:

The SocketChannel class reads from and writes to TCP sockets. The data must be encoded in ByteBuffer objects for reading and writing.

#### 1.1 Connecting:

The SocketChannel class does not have any public constructors. Instead, you create a new SocketChannel object using one of the two static open() methods:

***public static SocketChannel open(SocketAddress remote) throws IOException***

***public static SocketChannel open() throws IOException***

The first variant makes the connection. This method blocks (i.e., the method will not return until the

connection is made or an exception is thrown). For example:

```
SocketAddress address = new  
InetSocketAddress("www.cafeaulait.org", 80);  
SocketChannel channel = SocketChannel.open(address);
```

The noargs version does not immediately connect. It creates an initially unconnected socket that must be connected later using the `connect()` method. For example:

```
SocketChannel channel = SocketChannel.open();  
SocketAddress address = new  
InetSocketAddress("www.cafeaulait.org", 80);  
channel.connect(address);
```

## 1.2 Reading:

To read from a `SocketChannel`, first create a `ByteBuffer` the channel can store data in.

Then pass it to the `read()` method:

```
public abstract int read(ByteBuffer dst) throws IOException
```

## 1.3 Writing:

Socket channels have both read and write methods. In general, they are full duplex. In order to write, simply fill a `ByteBuffer`, flip it, and pass it to one of the write methods, which drains it while copying the data onto the output—pretty much the reverse of the reading process.

The basic `write()` method takes a single buffer as an argument:

```
public abstract int write(ByteBuffer src) throws IOException
```


## Scattering Reads and Gathering Writes

- Scattering Read: reads data from a single channel into multiple buffers

```
ByteBuffer tcpHeader = ByteBuffer.allocate(20);
```

```
ByteBuffer tcpBody = ByteBuffer.allocate(1220);
```

```
ByteBuffer[] bufferArray = { tcpHeader, tcpBody };  
channel.read(bufferArray);
```




- Gathering Write: writes data from multiple buffers into a single channel

```
ByteBuffer tcpHeader = ByteBuffer.allocate(20);
```

```
ByteBuffer tcpBody = ByteBuffer.allocate(1220);
```

```
//write data into buffers
```

```
ByteBuffer[] bufferArray = { tcpHeader, tcpBody };  
channel.write(bufferArray);
```





## 1.4 Closing:

Just as with regular sockets, you should close a channel when you're done with it to free up the port and any other resources it may be using:

*public void close() throws IOException*

## 2.ServerSocketChannel:

The ServerSocketChannel class has one purpose: **to accept incoming connections**. You cannot read from, write to, or connect a ServerSocketChannel. The only operation it supports is accepting a new incoming connection.

### 2.1 Creating server socket channels

The static factory method **ServerSocketChannel.open()** creates a new ServerSocketChannel object.

This method **does not actually open a new server socket**. Instead, it just **creates the object**.

Before you can use it, you need **to call the socket()** method to get the corresponding **peer ServerSocket**.

*try {*

```

ServerSocketChannel serverChannel =
ServerSocketChannel.open();

ServerSocket socket = serverChannel.socket();

SocketAddress address = new InetSocketAddress(80);

socket.bind(address);

} catch (IOException ex) {

System.err.println("Could not bind to port 80 because " +
ex.getMessage());

}

```

## 2.2 Accepting connections:

Once you've opened and bound a ServerSocketChannel object, the **accept()** method can **listen for incoming connections**:

```
public abstract SocketChannel accept() throws IOException
```

accept() can operate in either **blocking or nonblocking mode**.

## 3.The Channels Class:

Channels is a simple utility class for wrapping channels around traditional I/O-based streams, readers, and writers, and vice versa.

It has methods that convert from streams to channels and methods that convert from channels to streams, readers, and writers:

***public static InputStream newInputStream(ReadableByteChannel ch)***: Returns an input stream that reads bytes from the specified readable byte channel.

***public static OutputStream newOutputStream(WritableByteChannel ch)***: Returns an output stream that writes bytes to the specified writable byte channel.

***public static ReadableByteChannel newChannel(InputStream in)***: Returns a readable byte channel that reads bytes from the specified input stream.

***public static WritableByteChannel newChannel(OutputStream out)***: Returns a writable byte channel that writes bytes to the specified output stream.

***public static Reader newReader (ReadableByteChannel channel, CharsetDecoder decoder, int minimumBufferCapacity)***: Returns a reader that reads characters from the specified readable byte channel using the specified character decoder and minimum buffer capacity.

***public static Reader newReader (ReadableByteChannel ch, String encoding)***: Returns a reader that reads characters from the specified readable byte channel using the specified character encoding.

***public static Writer newWriter (WritableByteChannel ch, String encoding)***: Returns a writer that writes characters to the specified writable byte channel using the specified character encoding.

## 4. Asynchronous Channels

Java 7 introduces the `AsynchronousSocketChannel` and `AsynchronousServerSocketChannel` classes. These behave like and have almost the same interface as `SocketChannel` and `ServerSocketChannel`.

However, unlike `SocketChannel` and `ServerSocketChannel`, reads from and writes to asynchronous channels return immediately, even before the I/O is complete. The data read or written is further processed by a `Future` or a `CompletionHandler`. The `connect()` and `accept()` methods also execute asynchronously and return `Futures`. Selectors are not used.

For example, suppose a program needs to perform a lot of initialization at startup. Some of that involves network connections that are going to take several seconds each. You can start several asynchronous operations in parallel, then perform your local initializations, and then request the results of the network operations:

```
SocketAddress address = new InetSocketAddress(args[0], port);  
AsynchronousSocketChannel client =  
    AsynchronousSocketChannel.open();  
Future<Void> connected = client.connect(address);  
ByteBuffer buffer = ByteBuffer.allocate(74);
```

```

// wait for the connection to finish
connected.get();
// read from the connection
Future<Integer> future = client.read(buffer);
// do other things...
// wait for the read to finish...
future.get();
// flip and drain the buffer
buffer.flip();
WritableByteChannel out = Channels.newChannel(System.out);
out.write(buffer);

```

The advantage of this approach is that the network connections run in parallel while the program does other things. When you're ready to process the data from the network, but not before, you stop and wait for it by calling `Future.get()`.

## 5. Socket Options (Java 7)

Beginning in Java 7, `SocketChannel`, `ServerSocketChannel`, `AsynchronousServerSocketChannel`, `AsynchronousSocketChannel`, and `DatagramChannel` all implement the new `NetworkChannel` interface. The primary purpose of this interface is to support the various TCP options such as `TCP_NODELAY`, `SO_TIMEOUT`, `SO_LINGER`, `SO_SNDBUF`, `SO_RCVBUF`, and `SO_KEEPALIVE`.

the channel classes each have just three methods to get, set, and list the supported options:

*<T> T getOption(SocketOption<T> name) throws IOException*

-used to retrieve the value of the specified socket option

**Example:**

```
SocketChannel socketChannel = SocketChannel.open();
```

```
int timeout =
```

```
socketChannel.getOption(StandardSocketOptions.SO_TIMEOUT);
```

*<T> NetworkChannel setOption(SocketOption<T> name, T value)*

*throws IOException* : used to set the value of the specified socket option

**Example:**

```
SocketChannel socketChannel = SocketChannel.open();
```

```
socketChannel.setOption(StandardSocketOptions.SO_TIMEOUT,  
10000);
```

*Set<SocketOption<?>> supportedOptions()*:used to retrieve a set of all the socket options that are supported by the network channel.

**Example:**

```
SocketChannel socketChannel = SocketChannel.open();
```

```
Set<SocketOption<?>> options =
```

```
socketChannel.supportedOptions();
```

The StandardSocketOptions class provides constants for each of the 11 options Java recognizes:

- *SocketOption<NetworkInterface>*  
*StandardSocketOptions.IP\_MULTICAST\_IF*
- *SocketOption<Boolean>*  
*StandardSocketOptions.IP\_MULTICAST\_LOOP*
- *SocketOption<Integer>*  
*StandardSocketOptions.IP\_MULTICAST\_TTL*
- *SocketOption<Integer>* *StandardSocketOptions.IP\_TOS*
- *SocketOption<Boolean>*  
*StandardSocketOptions.SO\_BROADCAST*
- *SocketOption<Boolean>*  
*StandardSocketOptions.SO\_KEEPALIVE*
- *SocketOption<Integer>* *StandardSocketOptions.SO\_LINGER*
- *SocketOption<Integer>* *StandardSocketOptions.SO\_RCVBUF*
- *SocketOption<Boolean>*  
*StandardSocketOptions.SO\_REUSEADDR*
- *SocketOption<Integer>* *StandardSocketOptions.SO\_SNDBUF*
- *SocketOption<Boolean>* *StandardSocketOptions.TCP\_NODELAY*

### *Example 11-7. Listing supported options*

```
import java.io.*;
import java.net.*;
import java.nio.channels.*;
public class OptionSupport {
public static void main(String[] args) throws IOException {
    printOptions(SocketChannel.open());
    printOptions(ServerSocketChannel.open());
    printOptions(AsynchronousSocketChannel.open());
    printOptions(AsynchronousServerSocketChannel.open());
    printOptions(DatagramChannel.open());
}
private static void printOptions(NetworkChannel channel) throws
    IOException {
```

```
System.out.println(channel.getClass().getSimpleName() + "
supports:");
for (SocketOption<?> option : channel.supportedOptions()) {
    System.out.println(option.name() + ": " +
        channel.getOption(option));
}
System.out.println();
channel.close();
}
}
```

## Readiness Selection

The second part of the new I/O APIs is readiness selection, the ability to choose a socket that will not block when read or written

In order to perform readiness selection, different channels are registered with a Selector object. Each channel is assigned a SelectionKey.

The program can then ask the Selector object for the set of keys to the channels that are ready to perform the operation you want to perform without blocking.



## 1. The Selector Class

The only constructor in Selector is protected. Normally, a new selector is created by invoking the static factory method **Selector.open()**:

*public static Selector open() throws IOException*

The next step is to add channels to the selector. There are no methods in the Selector class to add a channel. the channel is registered with a selector by passing the selector to one of the channel's register methods.

*public final SelectionKey register(Selector sel, int ops)  
throws ClosedChannelException*

*public final SelectionKey register(Selector sel, int ops,  
Object att) throws ClosedChannelException*

The first argument is the selector the channel is registering with. The second argument is a named constant from the SelectionKey class identifying the operation the channel is registering for. The SelectionKey class defines four named bit constants used to select the type of the operation:

- SelectionKey.OP\_ACCEPT
- SelectionKey.OP\_CONNECT
- SelectionKey.OP\_READ

- `SelectionKey.OP_WRITE`

if a channel needs to register for multiple operations in the same selector, combine the constants with the bitwise or operator (`|`) when registering:

```
channel.register(selector, SelectionKey.OP_READ |  
SelectionKey.OP_WRITE);
```

The optional third argument is an attachment for the key

## 2. The `SelectionKey` Class

The `SelectionKey` class in Java is part of the Java NIO (New Input/Output) package, which provides a non-blocking I/O functionality.

It represents a key for a selection operation, allowing you to monitor the readiness of a channel for I/O operations, such as reading or writing data.

`SelectionKey` objects are returned by the `register()` method when registering a channel with a selector.

When retrieving a `SelectionKey` from the set of selected keys, you often first test what that key is ready to do. There are four possibilities:

```
public final boolean isAcceptable()  
public final boolean isConnectable()  
public final boolean isReadable()  
public final boolean isWritable()
```

Once you know what the channel associated with the key is ready to do, retrieve the channel with the `channel()` method:

```
public abstract SelectableChannel channel()
```

If you've stored an object in the `SelectionKey` to hold state information, you can retrieve it with the `attachment()` method:

```
public final Object attachment()
```

Finally, when you're finished with a connection, deregister its `SelectionKey` object so the selector doesn't waste any resources querying it for readiness.

```
public abstract void cancel()
```